

calamanCy: A Tagalog Natural Language Processing Toolkit

Lester James V. Miranda

ExplosionAI GmbH

lj@explosion.ai

Abstract

We introduce calamanCy, an open-source toolkit for constructing natural language processing (NLP) pipelines for Tagalog. It is built on top of spaCy, enabling easy experimentation and integration with other frameworks. calamanCy addresses the development gap by providing a consistent API for building NLP applications and offering efficient multitask models with out-of-the-box support for dependency parsing, part-of-speech (POS) tagging, and named entity recognition (NER). calamanCy aims to accelerate the progress of Tagalog NLP by consolidating disjointed resources in a unified framework. The calamanCy toolkit can be found on Github: <https://github.com/ljvmiranda921/calamanCy>.

1 Introduction

Tagalog is a low-resource language from the Austronesian family with over 76 million speakers in the Philippines (Lewis, 2009). Despite its speaker population, few resources exist for the language (Cruz and Cheng, 2021). For example, Universal Dependencies (UD) treebanks for Tagalog are tiny (\ll 20k words) (Samson, 2018; Aquino and de Leon, 2020), while domain-specific corpora are sparse (?). In addition, Tagalog language models (LMs) (Cruz and Cheng, 2021; Jiang et al., 2021) are few while most multilingual LMs (Conneau et al., 2019; Devlin et al., 2019) underrepresent the language. Thus, consolidating these disjointed resources in a coherent framework is still an open problem. The lack of such framework hampers model development, experimental workflows, and the overall advancement of Tagalog NLP.

To address this problem, we introduce calamanCy,¹ an open-source toolkit for Tagalog NLP. It is built on top of spaCy (Honnibal et al., 2020) and offers end-to-end pipelines for NLP tasks such

as dependency parsing, parts-of-speech (POS) tagging, and named entity recognition (NER). calamanCy also provides models of different sizes to fit any performance or accuracy requirements. This work has two main contributions: (1) an open-source toolkit containing multitask pipelines with out-of-the box support for common NLP tasks, and (2) the first structured benchmark that evaluates on several Tagalog core NLP tasks.

2 Related Work

Open-source toolkits for NLP There has been a growing body of work in the development of NLP toolkits in recent years. For languages, these toolkits include DaCy for Danish (Enevoldsen et al., 2021) and HuSpaCy for Hungarian (Orosz et al., 2022). For domain-specific data, there is medspaCy for clinical text (Eyre et al., 2021) and scispaCy for scientific articles (Neumann et al., 2019). These tools employ spaCy (Honnibal et al., 2020), an industrial-strength open-source software for natural language processing. Using spaCy as a foundation is an optimal choice given its popularity and tight integration with other frameworks such as HuggingFace (Wolf et al., 2019). However, no tool exists for Tagalog until now. In this paper, we will showcase how calamanCy provides similar capabilities to DaCy and HuSpaCy using Tagalog resources.

Evaluations on Tagalog NLP Tasks Structured evaluations for core NLP tasks, such as dependency parsing, POS tagging, and NER, are sparse. However, we have access to a reasonable amount of data to conduct comprehensive benchmarks. For example, TLUnified (Cruz and Cheng, 2021) is a pretraining corpus that combines news reports (Cruz et al., 2020), a preprocessed version of CommonCrawl (Suarez et al., 2019), and several other datasets. However, it was evaluated on domain-specific applications that may not easily transfer

¹The name “calamanCy” came from *kalamansi*, a citrus fruit native to the Philippines.

Entity	Description	Examples
Person (PER)	Person entities limited to humans. It may be a single individual or group.	Juan de la Cruz, Jose Rizal, Quijano de Manila
Organization (ORG)	Organization entities limited to corporations, agencies, and other groups of people defined by an organizational structure.	Meralco, DPWH, United Nations
Location (LOC)	Location entities are geographical regions, areas, and landmasses. Geo-political entities are also included within this group.	Pilipinas, Manila, CAL-ABARZON, Ilog Pasig

Table 1: Entity types used for annotating calamanCy-gold (derived from the TLUnified corpus of Cruz and Cheng, 2021). Annotation guidelines can be found at https://github.com/ljvmiranda921/calamanCy/tree/master/datasets/tl_calamancy_gold_corpus/guidelines

to more general tasks. For dependency parsing and POS tagging, we have Universal Dependencies treebanks such as TRG (Samson, 2018) and Ugnayan (Aquino and de Leon, 2020). This paper will fill the evaluation gap by providing structured benchmarks on these core tasks.

3 Implementation

The best way to use calamanCy is through its trained pipelines. After installing the library, users can access the models via:

```
import calamancy as cl
nlp = cl.load("tl_calamancy_md-0.1.0")
```

Here, the variable `nlp` is a spaCy processing pipeline.² It contains trained components for POS tagging, dependency parsing, and NER. calamanCy offers three pipelines of varying capacity: two static word vector-based models (`md`, `lg`), and one transformer-based model (`trf`). We will discuss how these pipelines were developed in the following section.

3.1 Pipeline development

Data annotation To construct the NER corpus, we curated a portion of TLUnified (Cruz and Cheng, 2021) to only contain Tagalog news articles. Including the author, we recruited two more annotators who have at least a Bachelors degree and whose native language is Tagalog. The three annotators labeled over the course of four months given three entity types as seen in Table 1. The entity types were chosen to resemble ConLL (Sang, 2002; Sang and Meulder, 2003), a standard NER benchmark. We measured inter-annotator agree-

Dataset	Examples	PER	ORG	LOC
Training	6252	6418	3121	3296
Development	782	793	392	409
Test	782	818	423	438

Table 2: Dataset statistics for calamanCy-gold.

ment (IAA) by taking the pairwise Cohen’s κ without the un-annotated tokens then averaged them for all three pairs. This process resulted to a Cohen’s κ score of 0.78. To avoid confusing with the original TLUnified corpora, we will refer to this annotated NER dataset as calamanCy-gold. The final dataset statistics can be found in Table 2. For the dependency parser and POS tagger, we merged the TRG (Samson, 2018) and Ugnayan (Aquino and de Leon, 2020) treebanks to leverage their small yet relevant examples.

Model training We considered three design dimensions when training the calamanCy pipelines: (1) presence of pretraining, (2) the word representation, and (3) the representation or dimension size. *Pretraining* involves learning vectors from raw text to better inform model initialization. This process is done using a variant of the cloze task (Devlin et al., 2019). Here, the pretraining objective asks the model to predict some number of leading and trailing UTF-8 bytes for the words. *Word representations* may either involve training static word embeddings using floret,³ an efficient version of fastText (Bojanowski et al., 2016), or using context-sensitive vectors from a transformer (Vaswani et al., 2017). Finally, the *dimension* is determined via a performance-accuracy tradeoff.

²<https://spacy.io/usage/processing-pipelines>

³<https://github.com/explosion/floret>

Pipeline	Pretraining objective	Word embeddings	Dimensions
Medium-sized pipeline (tl_calamancy_md)	Predict some number of leading and trailing UTF-8 bytes for the words.	Uses floret vectors trained on the TLUnified corpora.	50k unique vectors (200 dimensions), Size: 77 MB
Large-sized pipeline (tl_calamancy_lg)	Same pretraining objective as the medium-sized pipeline.	Uses fastText vectors trained on Common-Crawl corpora.	714k unique vectors (300 dimensions), Size: 455 MB
Transformer-based pipeline (tl_calamancy_trf)	No separate pretraining because there's no token-to-vector component.	Context-sensitive and vectors from a transformer network.	Uses roberta-tagalog-base. Size: 813 MB

Table 3: Language pipelines available in calamanCy (v0.1.0). The pretraining method for the word-vector models is a variant of the *cloze task*. All pipelines have a tagger, parser, morphologizer, and ner spaCy component.

The general process involves pretraining a filtered version of TLUnified, constructing static word embeddings if necessary, and training the downstream components. We trained the NER component using data from calamanCy-gold, while the dependency parser and POS tagger were trained using the combined TRG and Ugnayan treebanks. In the end, we came up with three language pipelines of varying sizes as seen in Table 3.⁴

4 Evaluation

Tasks and benchmark datasets We evaluated the calamanCy pipelines on various Tagalog benchmarks as seen in Table ?? . Unlike the *Hatespeech* and *Dengue* text categorization datasets, there is no reasonably-sized benchmark for NER and dependency parsing. So instead, we used a held-out test split from calamanCy-gold for NER and then merged the two UD treebanks. However, the combined UD treebank is still small ($\ll 20k$ words), so we evaluated it using 10-fold cross-validation.

Architectures

5 Discussions

6 Conclusion

References

Angelina A. Aquino and Franz A. de Leon. 2020. Parsing in the absence of related languages: Evaluating low-resource dependency parsers on Tagalog. In *Universal Dependencies Workshop*.

⁴The naming convention resembles spaCy's model names: {language code}_{source}_{size}. The full training configuration for the 0.1.0 version can be found on Github: <https://github.com/ljvimrainda921/calamanCy/tree/master/models/v0.1.0>

Piotr Bojanowski, Edouard Grave, Armand Joulin, and Tomas Mikolov. 2016. Enriching Word Vectors with Subword Information. *Transactions of the Association for Computational Linguistics*, 5:135–146.

Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Unsupervised Cross-lingual Representation Learning at Scale. In *Annual Meeting of the Association for Computational Linguistics*.

Jan Christian Blaise Cruz and Charibeth Ko Cheng. 2021. Improving Large-scale Language Models and Resources for Filipino. In *International Conference on Language Resources and Evaluation*.

Jan Christian Blaise Cruz, Jose Kristian Resabal, James Lin, Dan John Velasco, and Charibeth Ko Cheng. 2020. Exploiting News Article Structure for Automatic Corpus Generation of Entailment Datasets. In *Pacific Rim International Conference on Artificial Intelligence*.

Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *ArXiv*, abs/1810.04805.

Kenneth C. Enevoldsen, L M Hansen, and Kristoffer Laigaard Nielbo. 2021. DaCy: A Unified Framework for Danish NLP. In *Workshop on Computational Humanities Research*.

Hannah Eyre, Alec B. Chapman, Kelly S. Peterson, Jianlin Shi, Patrick R. Alba, Makoto M. Jones, Tamara L Box, Scott L Duvall, and Olga V Paterson. 2021. Launching into clinical space with medspaCy: a new clinical text processing toolkit in Python. *Proceedings of the AMIA Annual Symposium*, 2021:438–447.

Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. 2020. [spaCy: Industrial-strength Natural Language Processing in Python](#).

- Shengyi Jiang, Yingwen Fu, Xiaotian Lin, and Nankai Lin. 2021. Pre-trained Language Models for Tagalog with Multi-source Data. In *Natural Language Processing and Chinese Computing*.
- Paul M. A. Lewis. 2009. Ethnologue: languages of the world. <https://ethnologue.com/language/tgl>. Accessed: June 2023.
- Mark Neumann, Daniel King, Iz Beltagy, and Waleed Ammar. 2019. ScispaCy: Fast and Robust Models for Biomedical Natural Language Processing. *ArXiv*, abs/1902.07669.
- György Orosz, Zsolt Szántó, Péter Berkecz, Gergo Szabó, and Richárd Farkas. 2022. HuSpaCy: an industrial-strength Hungarian natural language processing toolkit. *ArXiv*, abs/2201.01956.
- Stephanie Dawn Samson. 2018. A treebank prototype of Tagalog. Bachelor’s thesis, University of Tübingen, Germany.
- Erik Tjong Kim Sang. 2002. Introduction to the CoNLL-2002 Shared Task: Language-Independent Named Entity Recognition. *ArXiv*, cs.CL/0209010.
- Erik Tjong Kim Sang and Fien De Meulder. 2003. Introduction to the CoNLL-2003 Shared Task: Language-Independent Named Entity Recognition. *ArXiv*, cs.CL/0306050.
- Pedro Ortiz Suarez, Benoît Sagot, and Laurent Romary. 2019. Asynchronous Pipeline for Processing Huge Corpora on Medium to Low Resource Infrastructures. In *7th Workshop on the Challenges in the Management of Large Corpora*.
- Ashish Vaswani, Noam M. Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. Attention is All you Need. In *Proceedings of the Conference on Neural Information Processing Systems*.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, and Jamie Brew. 2019. HuggingFace’s Transformers: State-of-the-art Natural Language Processing. *ArXiv*, abs/1910.03771.